

Securing Data-in-Transit with Calico & WireGuard

Wednesday, Sep 18, 2024

Prepared by Davide Sellitri,
Customer Success Engineer @ Tigera



Agenda



- Introduction
- The need for network encryption
- How it works
- Configuration Challenges
- Calico's Way
- Lab Description
- Demo

Introduction



Introduction: Calico Open Source



8M+
Nodes



1M+
Clusters



1.4B+
Docker Pulls



50k+
Enterprises



>50%
of Fortune 100



166
Countries


Most adopted container networking and security solution

Introduction: Calico Open Source



Introduction: WireGuard Tunneling Protocol

- simple, lightweight, and performant
- UDP
- kernel VNI
- CryptoKey Routing
- Cross Platform



The need for network encryption



Wireguard: How it works

- Provides **wg**, **wg-quick**, and **wgctrl-go** to configure Wireguard and view information
- Establishes a secure tunnel between two peers
 - Manual **configurations** on each peer
 - **CryptoKey Routing** to associate public keys with IP addresses & endpoint
 - Routing table when sending packets
 - Access control list when receiving packets


```
ubuntu@ip-10-0-1-30:~$ sudo wg
```

```
interface: wireguard.cali
```

```
public key: fpmFbBX4lMh7hH9vT4zgGnLypBcHz1XA81YFB7ZNAX8=
```

```
private key: (hidden)
```

```
listening port: 51820
```

```
fwmark: 0x20000000
```

```
peer: 0srB5mLCPfZKy19Fke/MAXuH0yjEt1Hq13e9ZfjX0Hk=
```

```
endpoint: 10.0.1.20:51820
```

```
allowed ips: 10.48.115.64/26, 10.48.115.64/32, 10.48.115.71/32
```

```
latest handshake: 15 seconds ago
```

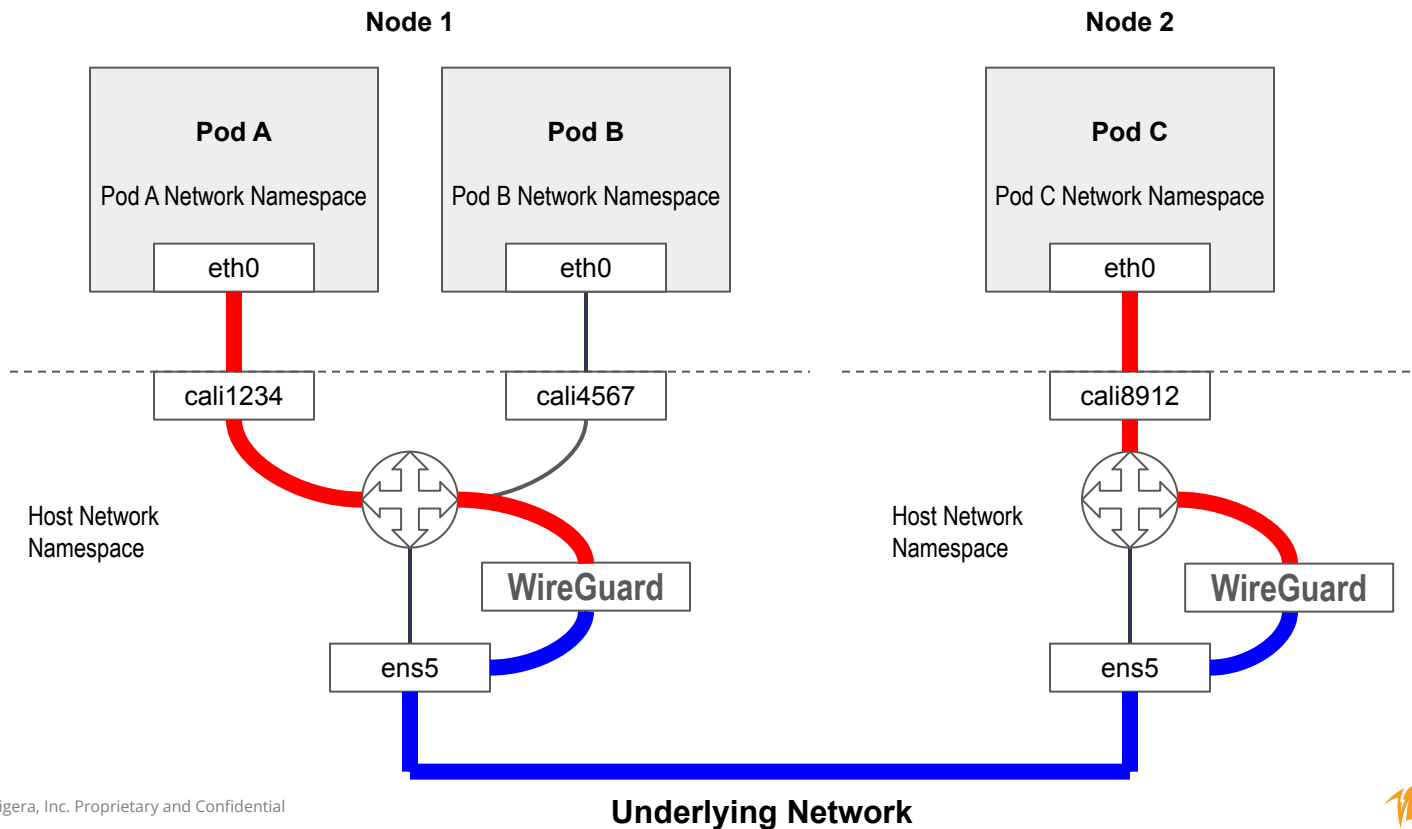
```
transfer: 17.88 KiB received, 154.35 KiB sent
```

```
peer: kw3cnxNy2CR/3c1LvVujhLpLI944rcyhkX2z6u3WuGk=
```

```
endpoint: 10.0.1.31:51820
```

```
allowed ips: 10.48.116.128/26, 10.48.116.128/32, 10.48.116.138/32
```

How it works: End Result



Wireguard Configuration Challenges

- A manual process
 - Create a key pair
 - Create tunnel virtual interface, allocate an IP address, and bring the interface online
 - Configure the endpoint address
 - Configure and continuously update the allowed IPs
 - No in-band key exchange: distribute public keys securely out-of-band
 - Keys are stored in memory and are lost upon node restarts

```
kubectl patch felixconfiguration default  
--type='merge' -p '{"spec":{"wireguardEnabled":true}}'
```

Lab Description

Bastion



Kubernetes Cluster

Worker 1



NC Client

Worker 2



NC Server

Port 12345

Demo



Q & A



GitHub



TIGERA

Follow us on:



Thank you!



Feedback



TIGERA

Follow us on:




Additional Slides




Calico Wireguard Widget in Calico Cloud or Calico Enterprise

Wireguard. Handshake



1.85

Wireguard. Bytes sent/received



157k↓
157k↑
Per Sec

Calico – Wireguard Encryption Support

- Provides in-cluster pod traffic encryption
- Supported platforms:

The following platforms using only IPv4:

- Kubernetes, on-premises
- EKS using Calico CNI
- EKS using AWS CNI
- AKS using Azure CNI
- AKS using Calico CNI

All platforms listed above will encrypt pod-to-pod traffic. Additionally, when using AKS or EKS, host-to-host traffic will also be encrypted, including host-networked pods.

Implementation ...

- Only on nodes that have WireGuard installed
- Rules from local workload to the wireguard table

```
$ ip rule
0:    from all lookup local
99:   from all fwmark 0x0/0x800000 lookup 1
32766: from all lookup main
32767: from all lookup default
```

- FWMark is used to prevent routing loops
 - First pass routes packets to wireguard (in table 1)
 - Second pass routes wireguard encrypted UDP packets to the other host (in main table)

Implementation ...

- There is a new wireguard route table (index 1)



```
$ ip route show table 1  
throw 192.168.6.128/26  
192.168.24.128/26 dev wireguard.cali scope link  
192.168.42.64/26 dev wireguard.cali scope link
```

- Routes to pods on hosts that have wireguard installed via wireguard interface
- Routes to local workloads, or to pods on hosts that do not have wireguard installed revert back to previous routing (throw routes)

Implementation ...

- Supported config options in FelixConfig:
 - wireguardInterfaceName
 - wireguardListeningPort
 - wireguardMTU
 - wireguardRoutingRulePriority
 - wireguardKeepAlive

Wireguard Topologies



Wireguard Configuration – Calico's Way

- A felix configuration parameter
- Uses **wgctrl-go** to configure Wireguard
- Configures a wireguard virtual interface named **wireguard.cali**
- Key pairs are stored in memory and regenerated upon the node restart
- Stores public key in the datastore and securely distribute them between the nodes
- Uses the default Wireguard port 51820
- Configures the peers by accessing the required information in the datastore
 - Public key
 - Endpoint address – node API resource
 - Allowed IPs – Pod IP addresses and IPAM blocks
- Configures ip rule and special routing table
- Configures iptables/ebpf rules to allow Wireguard traffic
- Encrypts pods traffic across different nodes
 - Intra-node pod connectivity is not encrypted by Wireguard

Wireguard Configuration

- Load Wireguard kernel module

1

Peer1

2

```
wg genkey | sudo tee /etc/wireguard/private.key
sudo cat /etc/wireguard/private.key | wg pubkey | sudo tee /etc/wireguard/public.key
sudo ip link add dev wg0 type wireguard
sudo ip address add dev wg0 10.16.1.220/24
sudo ip link set up dev wg0
```

Peer2

2

```
wg genkey | sudo tee /etc/wireguard/private.key
sudo cat /etc/wireguard/private.key | wg pubkey | sudo tee /etc/wireguard/public.key
sudo ip link add dev wg0 type wireguard
sudo ip address add dev wg0 10.16.1.221/24
sudo ip link set up dev wg0
```

Wireguard Configuration

```
root@peer1:~#vi /etc/wireguard/wg0.conf
```

3

[Interface]

```
PrivateKey = yB8I/xYvMRCE5pjb7CFv3zbespTqNLNcX0FR0S1FmFg=  
ListenPort = 51820
```

[Peer]

```
PublicKey = q9uKGNlQOSwAD+/4EA+9A1lp1lqyHtEN/qOyCf8mGms=  
Endpoint = 172.16.1.221:51820  
AllowedIPs = 10.16.1.221/32
```

4

```
sudo wg setconf wg0 wg0.conf
```

```
root@peer2:~#vi /etc/wireguard/wg0.conf
```

3

[Interface]

```
PrivateKey = GJzf73I/fdKmIAozFF9Qd1LGLBG52HWyj2zIAVgogG0=  
ListenPort = 51820
```

[Peer]

```
PublicKey = lNFJiC37+WwGAYK8LjJ5PGB+LbraoRlwOlOdCtS3Rk4=  
Endpoint = 172.16.1.220:51820  
AllowedIPs = 10.16.1.220/32
```

4

```
sudo wg setconf wg0 wg0.conf
```

Troubleshooting

- Tcpcdump on the wireguard interface (wireguard.cali)
 - Should see packets routed to wireguard and decrypted response
- Tcpcdump on the main host interface
 - Should see UDP packets between the wireguard supported nodes
- Wireguard shares public key information through the Node resource status
 - Check that the node has a public key assigned
 - If wireguard is enabled and installed, but there is no public key, check Felix logs.

```
metadata:  
  annotations:  
    projectcalico.org/WireguardPublicKey:  
      j1kVYyQYooZYzI2wFfNhSZez5eWh44yfq1wKVjLvSX  
      gY=
```

Wireguard Configuration – Enable Wireguard

- `kubectl patch felixconfiguration default --type='merge' -p '{"spec":{"wireguardEnabled":true}}'`
- `kubectl patch felixconfiguration default --type='merge' -p '{"spec":{"wireguardEnabledV6":true}}'`
- `kubectl patch felixconfiguration default --type='merge' -p '{"spec":{"wireguardEnabled":true,"wireguardEnabledV6":true}}'`

Wireguard Configuration Validation

```
# kubectl get nodes -o yaml | grep -i wireguard
```

```
    projectcalico.org/IPv4WireguardInterfaceAddr: 192.168.0.106  
    projectcalico.org/WireguardPublicKey:  
G/syC2zarrAsAJp6iF9+g73ym+m8Wb+limuqpnqoDWI=
```

```
    projectcalico.org/IPv4WireguardInterfaceAddr: 192.168.1.232  
    projectcalico.org/WireguardPublicKey:  
3hwASDwR/miQ6JLB53j2xx+3ZbqJsVtW0OLj0mb8+ns=
```

```
    projectcalico.org/IPv4WireguardInterfaceAddr: 192.168.1.108  
    projectcalico.org/WireguardPublicKey:  
Vz60TLji5oRN46EMHrU649TZj07fzGLhT3FIgfWd0XA=
```

IP Route Configurations

```
# ip rule list lookup 1
```

```
99:      not from all fwmark 0x2000000/0x2000000 lookup 1
```

```
----
```

```
# ip route list table 1
```

```
throw 192.168.0.96/28
```

```
192.168.1.96 dev wireguard.cali scope link
```

```
192.168.1.96/28 dev wireguard.cali scope link
```

```
192.168.1.108 dev wireguard.cali scope link
```

```
192.168.1.112/28 dev wireguard.cali scope link
```

```
192.168.1.208 dev wireguard.cali scope link
```

```
192.168.1.208/28 dev wireguard.cali scope link
```

```
192.168.1.224/28 dev wireguard.cali scope link
```

```
30 192.168.1.232 dev wireguard.cali scope link
```



TIGERA

IPTables Configuration

```
# iptables-save | grep -i wireguard
```

```
-A cali-INPUT -p udp -m comment --comment "cali:w1xlp5qwD97H6wH1" -m  
comment --comment "Allow incoming IPv4 Wireguard packets" -m multiport  
--dports 51820 -m addrtype --dst-type LOCAL -j ACCEPT
```

```
-A cali-OUTPUT -p udp -m comment --comment "cali:XjqWdMh-cno69n2T" -m  
comment --comment "Allow outgoing IPv4 Wireguard packets" -m multiport  
--dports 51820 -m addrtype --src-type LOCAL -j ACCEPT
```

```
-A cali-POSTROUTING -o wireguard.cali -m comment --comment  
"cali:kgfCOPW4UKtzMAmO" -m addrtype ! --src-type LOCAL --limit-iface-out  
-m addrtype --src-type LOCAL -j MASQUERADE --random-fully
```